

Respond to:#104, EuroSys 2021, "FirmGuide: Re-hosting the Linux Kernel of Embedded Firmware through Model-Guided Kernel Execution"

Contents

1 Introduction	1
2 Reviewer B	1
3 Reviewer F	3
A An example of model templates	3

1 Introduction

Dear Reviewers,

We would like to sincerely thank you for providing us positive and inspiring comments that will significantly improve the paper's quality. In this extended response letter, we will answer the concerns that are not detailed in the short response, especially technical details. Reviewers' comments are in the numbered italic face, and our responses are in the following. The responses to Reviewer A, C, D, and E are absent because we think our short response letter has already covered their questions.

2 Reviewer B

Comment 1: *Section 2.1: "Whereas the processor can access RAM directly through the system bus, I/O peripherals are accessed ... through the Memory Mapped I/O (MMIO) mechanism." – As far as I'm aware, all MMIO reads and writes go directly through the system bus just like access RAM (which just sits at a different part of the address space).*

Response: From software's perspective, accessing MMIO regions are similar as accessing memorys. However, from QEMU's perspective, it needs to call the read/write callbacks to emulate the hardware's logic.

Comment 2: *Section 3.3: This motivating example is confusing. In Figure 3a, the IRQ mask function takes in the IRQ number as a parameter to the function. However, you treat this as a fixed value (3) in your example. Why can't this be a symbolic value itself?*

TODO

Response: It can be a symbolic value. The semantics of `irq_mask_callback` is inferred in our device model. Thus we know we should put an IRQ number as a parameter. Thus this is an optimization to avoid constraint solve being stuck and this concrete value can be one of our execution goals.

Comment 3: *Also, these #define'd values (e.g., `INTC_REG_MASK`) are not usually the MMIO register physical addresses themselves, but rather an index into the device's MMIO address region as specified in the device tree (which are then used as an offset for the `ioremap()`'d region).*

Response: We use the physical address in this example for simplicity. In practice, we hooked the MMIO read/write functions to capture the real physical address kernel accesses for a specific device. Therefore, it doesn't matter whether the macro is defined as offsets or addresses.

Comment 4: *Section 4.1: Could you talk a little bit more of the format for the template? I know in the evaluation you cite some number of lines of C code for the two template models (interrupt controller and timer), but I have no sense for what these might actually look like. Are the transition condition functions linked in later based on the model parameters?*

TODO

Response: The format for the template is implemented in a generic way. Transition condition functions are treated as configurations that should be registered into the template. We provide an example of the template in Appendix [A](#).

Comment 5: *Section 4.2: Under "Parameter Inference", what happens if there is no prior value for the `<MMIOR, addr, USE_LAST_VALUE>` R/W seq node? Is this going to use some fixed initial value or?*

Response: In this case, we use fixed initial value 0.

Comment 6: *Section 4: For the timer, is the notion of passing time captured in your manually-generated model?*

Response: Yes. The kernel will always first generate a time value in the unit of `cycle_t`. Then it will convert that value to a MMIO value (usually some bit shift operations like `y = x << 2`). FirmGuide first collects the conversion expression during the symbolic execution, and then tries to understand when the kernel wants the next time interrupt by converting the MMIO value back to `cycle_t` according to the inverse function of the expression (or constraint solver).

Comment 7: *Table 1: You note that the timer is "Not Necessary" for the first two subtargets, because it is actually already supported in QEMU. Why did you not generate the device model anyway and use it instead of the one provided by QEMU?*

Response: In the two "Not Necessary" cases, the QEMU MIPS CPU emulation code already contains these device's implementation and we cannot replace or disable them easily. The large difference between paths and solutions is caused by the difference in SoC's kernel code. One path means a possible boot path of an SoC and usually corresponds to one solution.

3 Reviewer F

Comment 1: *"During this process, for each MMIO read operation, ...": this really needs an illustration to make sense.*

Response: During the simplified boot process, we hook MMIO read/write functions. For each MMIO read, we tell KLEE to introduce a new read-from symbol. That's because the MMIO read operation is indeed a volatile memory operation that has no assumption to keep its value all the time.

Comment 2: *"After the execution ...": are you talking about concrete or symbolic execution here? Also, what is it you're executing here?*

Response: This execution means symbolic execution. We compile and link kernel's code to an LLVM IR file. Then we use KLEE to do symbolic execution on this IR file. We also add our main function into this IR file to control the code KLEE executes. KLEE will execute the kernel's boot process (`start_kernel` function). And during the KLEE's execution, the only place introduces new symbols is the MMIO read operation.

Comment 3: *"we may get several paths that can successfully finish the whole booting process": how do you measure success here?*

Response: A successful path must reach the end of the booting process without any error or early exit during the symbolic execution.

Comment 4: *"... are modeled as dummy MMIO memory space that accepts read and write operations.": What does accepting read/write operations mean?*

Response: It means that Type-II device's MMIO regions are emulated as RAM memory regions. No hardware specific logic is emulated and kernel could read/write it as normal memory.

Comment 5: *"Kernel first determines a value in cycle_t ... Then it aligns the unit with the specific timer device.": What does "aligning the unit" mean?*

Response: Linux kernel and a real timer device use different units for time. Linux kernel uses the second(s) of which we are aware. The timer device use their smallest cycle (`cycle_t`) as unit, for example, $1\text{M cycle_t} = 1\text{s}$. The Linux kernel accepts a command from us, say setting a timer in 1 second, then the Linux kernel will first convert (align) 1 s to 100M `cycle_t`. Once kernel wants to set a countdown value to a specific Timer device, it will always first generate a time value in the unit of `cycle_t`. Then kernel needs to convert that value to the value set to Timer's MMIO region which has equivalent meaning (usually the conversion are some bit shift operations like $y = x \ll 2$). For example, if kernel wants to set 1s (100 `cycle_t`), it will set 25 (according to the example $y = x \ll 2$) to the Timer. The "aligning the unit" represents for this unit conversion process.

A An example of model templates

```
1 // interrupt_controller_template.c
2 // This is just a demo case that is not runnable.
3 static void dispatch_mmio_rw(
4     void *opaque, MMIO_DIR mmio_dir,
```

```

5     hwaddr offset, unsigned size, uint64_t value){
6     if (EVENT event = rw_seq_matched(
7         opaque, mmio_dir, offset, size, value))
8         STATE state = transit_state(opaque, event);
9         run_state_action(opaque, state, event);
10 }
11
12 static uint64_4 intc_read(
13     void *opaque, hwaddr offset, unsigned size){
14     // ...
15     dispatch_mmio_rw(opaque, MMIO_READ, offset, size, NULL);
16     // ...
17
18 static void intc_write(
19     void *opaque, hwaddr offset, uint64_t value, unsigned size){
20     // ...
21     dispatch_mmio_rw(opaque, MMIO_WRITE, offset, size, value);
22     // ...
23 }
24
25 static const MemoryRegionOps intc_ops = {
26     .read = intc_read,
27     .write = intc_write,
28 }
29
30 static void intc_init(Object *obj) {
31     CONFIG config = get_registered_config();
32     s->config = config;
33     memory_region_init_io(&s->mmio, ..., &intc_ops, ..., conf->mmio_len);
34     s->state = STATE_IDLE;
35 }
36
37 static TypeInfo intc_type_info = {
38     // ...
39     .instance_init = intc_init,
40     // ...
41 };
42
43 static void intc_types(void) {
44     type_register_static(&intc_type_info);
45 }
46
47 type_init(intc_types)

```